


Verifier-Guided GRPO — training pipeline & verifiable evaluation


One RDKit symbolic verifier: the **reward** during training, the **exact scorer** at evaluation.

TRAINING PIPELINE



↓ base model and trained checkpoint are evaluated under the **identical** protocol

EVALUATION · MOLECULARIQ BENCH



Why it's a clean RL signal: the same symbolic verifier scores every completion exactly, so the training reward and the benchmark are the same oracle — no learned reward model, no LLM judge, and no train/test leakage. pass@k then separates RL **sharpening** from genuine capability **expansion**.