


Verifier-Guided GRPO — pipeline

One RDKit verifier: reward at train time, exact scorer at eval time.



TRAINING



repeated each batch — reward variance across the G samples drives the gradient

EVALUATION · MOLECULARIQ BENCH



Same symbolic verifier for **reward** and **benchmark** — no learned reward model, no LLM judge, no leakage.